

Unigornel

Initializing Go in Mini-OS (Part 2)

Maxim Bonnaerens

March 7, 2016

Recap - Let's crash it!¹

Breakpoint 1, [...] at 0x6505f: x_cgo_sys_thread_create
(gdb) disas

Dump [...] for function x_cgo_sys_thread_create:

```
0x65050 <+0>:  sub    $0x18,%rsp
0x65054 <+4>:  mov     %rdi,%rdx
0x65057 <+7>:  mov     %rsi,%rcx
0x6505a <+10>: mov     %rsp,%rdi
0x6505d <+13>: xor     %esi,%esi
=> 0x6505f <+15>: mov     %fs:0x28,%rax
0x65068 <+24>: mov     %rax,0x8(%rsp)
0x6506d <+29>: xor     %eax,%eax
0x6506f <+31>: callq   0x9f3f <pthread_create>
0x65074 <+36>: test    %eax,%eax
0x65076 <+38>: jne     0x6508d
                                <x_cgo_sys_thread_create+61>
```

¹On Ubuntu Wily with gcc 5.2.1/4.9.3

Recap - Let's crash it!²

Breakpoint 1, [...] at 0x64fc0: x_cgo_sys_thread_create
(gdb) disas

Dump [...] for function x_cgo_sys_thread_create:

```
=> 0x64fc0 <+0>:  sub    $0x18,%rsp
    0x64fc4 <+4>:  mov     %rdi,%rdx
    0x64fc7 <+7>:  mov     %rsi,%rcx
    0x64fca <+10>: lea     0x8(%rsp),%rdi
    0x64fcf <+15>:  xor     %esi,%esi
    0x64fd1 <+17>:  callq   0x9f13 <pthread_create>
    0x64fd6 <+22>:  test    %eax,%eax
    0x64fd8 <+24>:  jne     0x64fdf
                                <x_cgo_sys_thread_create+31>
```

²On Debian Jessie with gcc 4.9.2

Fix crashes

- ▶ Crash on Ubuntu: Using uninitialized %fs-segment
- ▶ Crash on Debian: Stack cookies
 - ▶ edit src/cmd/go/build.go:

```
case "netbsd":  
    a = append(a, "-fno-stack-protector")  
}
```

Go Runtime Initialization - Implement GDT

//seg_desc_t represents one entry in the GDT.

```
struct _segment_descriptor {
    unsigned lolimit:16;    /* low bits of maximum segment
    unsigned lobase:24;      /* low bits of segment starting
    unsigned accessed:1;    /* segment accessed */
    unsigned rw:1;          /* readable for code, writable
    unsigned dc:1;          /* direction for data, conforming
    unsigned ex:1;          /* executable for code */
    unsigned system:1;      /* system critical descriptor */
    unsigned dpl:2;         /* descriptor privilege level */
    unsigned present:1;     /* present bit, must be 1 */
    unsigned hilimit:4;     /* high bits of maximum segment
    unsigned zero:2;        /* always zero */
    unsigned size:1;        /* size of p-mode segment */
    unsigned gran:1;        /* granularity of segment 1b for
    unsigned hibase:8;      /* high bits of maximum segment
} __attribute__((__packed__));
```

Implementing Global Descriptor Table - Initialize GDT

```
void init_gdt(void) {
    pte_t pte; unsigned long frames[1];

    seg_desc_fill(&gdt[SEG_DESC_CS], seg_desc_type_era);
    seg_desc_fill(&gdt[SEG_DESC_DS], seg_desc_type_rwa);
    seg_desc_fill(&gdt[SEG_DESC_FS], seg_desc_type_rwa);

    pte = __pte((virt_to_mach(&gdt)) | L1_PROT_RO);
    HYPERVISOR_update_va_mapping((unsigned long)&gdt,
        pte, UVMF_INVLPG)

    frames[0] = virt_to_mfn(&gdt);
    HYPERVISOR_set_gdt(frames, NUM_GDT_ENTRIES)

    memcpy(&fs, &gdt[SEG_DESC_FS], sizeof(fs));
    fs_pa = virt_to_mach(&gdt[SEG_DESC_FS]);
}
```

Implementing Global Descriptor Table - Initialize GDT

- ▶ Initialize GDT when kernel boots
 - ▶ Call `init_gdt()` in `kernel.c`

```
void start_kernel(void) {  
    [...]  
  
    /* Init memory management. */  
    init_mm();  
  
    /* Init GDT */  
    init_gdt();  
  
    [...]  
}
```

Implementing Global Descriptor Table - Update fs

```
#define KERNEL_FS ((3 << 3) | 1)

void switch_fs(unsigned long p) {
    fs.desc.lobase = p & 0xFFFFFFFF;
    fs.desc.hibase = (p >> 24) & 0xFF;
    HYPERVISOR_update_descriptor(fs_pa, fs.raw)

    asm volatile("mov %0, %%fs" :: "r"(KERNEL_FS));
}
```


Using fs-segment without crashing

```
//- go_main.c
void thread_main(void *ctx) {
    [...]
    tls_page = alloc_page();
    switch_fs(tls_page + PAGE_SIZE);

    /* TEST TLS */
    *(int *) (tls_page + PAGE_SIZE - 8) = 9876;
    asm volatile(
        "mov %%fs:0xfffffffffffffff8, %0" : "=r"(v)
    );

    if(v != 9876) {
        printk("Could not successfully set up TLS\n");
        *(char *)0 = 0;
    }
    [...]
}
```

Stubbed pthread functions give trouble

- ▶ Runtime initialization calls `pthread_attr_getstacksize()`³
 - ▶ Stubbed
 - ▶ Runtime thinks stacksize is 0
 - ▶ Calls functions for more stack \implies crashes
- ▶ Solution:
 - ▶ Return `STACK_SIZE` used by Mini-OS scheduler⁴

³`go/src/runtime/cgo/gcc_netbsd_amd64.c`

⁴`minios/arch/x86/sched.c`

Table of Contents

Go Runtime Initialization

- Recap

- Implementing GDT

- Using fs-segment without crashing

- Stubbed pthread functions give trouble