

Unigornel

Initializing Go in Mini-OS (Part 4)

Henri Verroken

March 15, 2016

Recap - Remarks on GDT Implementation

- ▶ Why initialize CS and DS, but not SS?
 - ▶ CS and DS segment descriptors should not be in our GDT
 - ▶ CS, DS and SS registers point to GDT entries added by Xen¹
 - ▶ ES is not used
- ▶ What about 64-bit base addresses?
 - ▶ Use MSR for upper 32-bits
 - ▶ We limit ourselves to addresses below 0xFFFFFFFF

¹In Mini-OS: `include/xen/arch-x86/xen-x86-64.h`

Recap - Remarks on GDT Implementation²

```
[...]
-#define KERNEL_CS ((1 << 3) | 1)
-#define KERNEL_DS ((2 << 3) | 1)
-#define KERNEL_FS ((3 << 3) | 1)
-
-#define SEG_DESC_CS (KERNEL_CS >> 3)
-#define SEG_DESC_DS (KERNEL_DS >> 3)

+#define KERNEL_FS ((1 << 3) | 1)
#define SEG_DESC_FS (KERNEL_FS >> 3)
[...]
```

```
void init_gdt(void) {
    [...]
-    seg_desc_fill(&gdt[SEG_DESC_CS], seg_desc_type_era);
-    seg_desc_fill(&gdt[SEG_DESC_DS], seg_desc_type_rwa);
    seg_desc_fill(&gdt[SEG_DESC_FS], seg_desc_type_rwa);
    [...]
}
```

²<https://github.ugent.be/unigornel/minios/commit/614dbe70b3968da67d645e23bd5e8ec059fc17dc>

Recap - Remarks on GDT Implementation³

```
void switch_fs(unsigned long p) {  
+   ASSERT((p & 0xFFFFFFFF) == p,  
+           "FS base is limited to 32 bits");  
  
    fs.desc.lobase = p & 0xFFFFFFFF;  
    fs.desc.hibase = (p >> 24) & 0xFF;  
    HYPERVISOR_update_descriptor(fs_pa, fs.raw);  
  
    __asm__ __volatile__(  
        "mov %0, %%fs" :: "r"(KERNEL_FS)  
    );  
}
```

³<https://github.ugent.be/unigornel/minios/commit/91817cfa8e99299355d220f8f111f02e8fae4c65>

Recap - Go Runtime Memory Requirements

- ▶ Mini-OS
 - ▶ Pseudo-physical = virtual
 - ▶ Page tables filled at startup and never changed
 - ▶ Small virtual address range
 - ▶ Cannot use high addresses
- ▶ Go Runtime
 - ▶ Uses `mmap` with `PROT_NONE` to reserve around 500GB
 - ▶ Allocates from that region when needed
 - ▶ Uses high addresses
- ▶ Mini-OS and Go are incompatible
 - ▶ Try to limit memory requirements in Go

Memory

- ▶ Decrease memory requirements in Go⁴

```
const (  
    _MHeapMap_TotalBits = 26 // 64 MByte  
    _MaxMem = uintpr(1 << _MHeapMap_TotalBits - 1)  
)  
  
arenaSize := round(_MaxMem, _PageSize)
```

- ▶ Increase VM memory to 256 MB
- ▶ Implement basic `mmap` in Mini-OS⁵

⁴<https://github.ugent.be/unigornel/go/commit/d37f5677be6a604c41e713ed3d186fc110ce1a06>

⁵<https://github.ugent.be/unigornel/minios/commit/fcde24adc6db56e43cf48c1e46106b6151661ffc>

Synchronization

- ▶ Go implements locks using OS primitives
 - ▶ On NetBSD using semaphores⁶
 - ▶ Platform-specific `semasleep` and `semawake`⁷
 - ▶ Using NetBSD `lwp_park` and `lwp_unpark`
 - ▶ `lwp` has quite difficult specs
 - ▶ On Linux using `futex`-mechanism⁸
 - ▶ `futex` has a very easy specification
- ▶ Implement basic `futex` wait/wake in Mini-OS⁹
- ▶ Change Go to use `futexes` for NetBSD¹⁰

⁶`ln runtime/lock_sema.go`

⁷`ln runtime/os1_netbsd.go`

⁸`ln runtime/lock_futex.go`

⁹<https://github.ugent.be/unigornel/minios/commit/eeacb0f2659ea8dddc155360a15cf3188f2abc5d>

¹⁰<https://github.ugent.be/unigornel/go/commit/b3d174d7c53ec33327beaf709b70889cf7e56830>

Various System Calls

- ▶ Implement some more system calls
 - ▶ `thread_id` for thread identification
 - ▶ Remove all `signal` stuff
 - ▶ Use valid `argc`, `argv` and `envs`
 - ▶ Implement `usleep`
 - ▶ Implement `now`
 - ▶ Implement `write` to print to console
 - ▶ Avoid open-ing `/dev/random` for random data^{11,12}

¹¹The current implementation of `getRandomData` uses only `nanotime` as its source of entropy.

¹²<https://github.ugent.be/unigornel/go/commit/6cbdecf80df6a62516faccf0ff2b966d7dd5370e>

Deadlock - Stuck in the GC

- ▶ Stuck in the GC thread
- ▶ All other threads were blocking¹³
 - ▶ Fix memory bug in futex-implementation
 - ▶ Implement variant of pthreads conditions

¹³<https://github.ugent.be/unigornel/minios/commit/f7a9121e29cbf108b4a188c06060cdd6a365df5c>

Hello World - Hello Sum^{14,15}

```
package main

import "C"

func main() {}

//export Sum
func Sum(a, b int) int {
    return a + b
}
```

¹⁴Mini-OS at f7a9121e29cbf108b4a188c06060cdd6a365df5c

¹⁵Go at d7d2d60a91bd0cb79b09e07a4468c0fcf4c7d711

Hello World - Hello Sum^{16,17}

```
static void *initialize_go_thread(void *ctx) {
    _rt0_amd64_netbsd_lib();
    return NULL;
}

static void *main_thread(void *ctx) {
    GoInt i = Sum(3, 4);
    printk("main_thread: result: %lld\n", i);
    CRASH("main thread must not return");
    return NULL;
}

int app_main(start_info_t *si) {
    pthread_t t;
    pthread_create_name(&t, "main", NULL,
                       main_thread, NULL);
    pthread_create_name(&t, "initialize_go", NULL,
                       initialize_go_thread, NULL);

    return 0;
}
```

¹⁶Mini-OS at f7a9121e29cbf108b4a188c06060cdd6a365df5c

¹⁷Go at d7d2d60a91bd0cb79b09e07a4468c0fcf4c7d711

Hello World - Hello Sum

```
[...]  
go_main.c: app_main(0xffe40)  
Thread "main": pointer: 0x2b2220, stack: 0x2f0000  
Thread "initialize_go": pointer: 0x2b2300, stack: 0x300000  
mallocinit: arenaSize 64 MByte  
mmap(addr=[...],len=0x4412000,prot=0x0, flags=0x1002)  
mheap_.spans = 0x310000  
mheap_.bitmap = 0x320000  
mheap_.arena_start = 0x720000  
mheap_.arena_used = 0x720000  
mheap_.arena_end = 0x4722000  
mheap_.arena_reserved = true  
mmap(addr=0x0,len=0x40000,prot=0x3, flags=0x1002)  
mmap(addr=0x720000,len=0x100000,prot=0x3, flags=0x1002)  
mmap(addr=0x718000,len=0x8000,prot=0x3, flags=0x1002)  
mmap(addr=0x310000,len=0x1000,prot=0x3, flags=0x1002)  
mmap(addr=0x0,len=0x10000,prot=0x3, flags=0x1002)  
Thread "pthread-0": pointer: 02b2488, stack: 04780000  
mmap(addr=0x0,len=0x40000,prot=0x3, flags=0x1002)  
Thread "pthread-1": pointer: 02b25a0, stack: 047d0000  
main_thread: result: 7  
crash: go_main.c, line 25: main thread must not return
```

Hello World^{18,19}

```
package main

import "C"
import "fmt"

func main() {}

//export Main
func Main(unused int) {
    fmt.Println("Hello World!")
}
```

¹⁸Mini-OS at c1b830d41e4612c49c95bb4c8b906137361a92c6

¹⁹Go at 6cbdecf80df6a62516faccf0ff2b966d7dd5370e

Hello World^{20,21}

```
static void *initialize_go_thread(void *ctx) {
    _rt0_amd64_netbsd_lib();
    return NULL;
}

static void *main_thread(void *ctx) {
    Main(0);
    CRASH("main thread must not return");
    return NULL;
}

int app_main(start_info_t *si) {
    pthread_t t;
    pthread_create_name(&t, "main", NULL,
                       main_thread, NULL);
    pthread_create_name(&t, "initialize_go", NULL,
                       initialize_go_thread, NULL);

    return 0;
}
```

²⁰Mini-OS at c1b830d41e4612c49c95bb4c8b906137361a92c6

²¹Go at 6cbdecf80df6a62516faccf0ff2b966d7dd5370e

Hello World

```
[...]
mheap_.spans = 0x380000
mheap_.bitmap = 0x390000
mheap_.arena_start = 0x790000
mheap_.arena_used = 0x790000
mheap_.arena_end = 0x4792000
mheap_.arena_reserved = true
mmap(addr=[...],len=0x40000,prot=0x3, flags=0x1002)
mmap(addr=[...],len=0x100000,prot=0x3, flags=0x1002)
mmap(addr=[...],len=0x8000,prot=0x3, flags=0x1002)
mmap(addr=[...],len=0x1000,prot=0x3, flags=0x1002)
mmap(addr=[...],len=0x10000,prot=0x3, flags=0x1002)
Thread "pthread-0": pointer: 0x329488, stack: 0x47f0000
mmap(addr=[...],len=0x40000,prot=0x3, flags=0x1002)
Thread "pthread-1": pointer: 0x3295a0, stack: 0x4840000
Hello World!
crash: go_main.c, line 21: main thread must not return
```

What's Next

- ▶ Clean up code
- ▶ Rewrite parts in Go (e.g. `futex`)
- ▶ Test with more complex applications
 - ▶ Do some memory intensive things
 - ▶ Use channels and multiple goroutines
- ▶ Support full virtual address space
- ▶ Explore networking in Xen
- ▶ Automated testing

Table of Contents

Recap

GDT

Go Runtime Memory Requirements

Memory

Synchronization

Various System Calls

Deadlock - Stuck in the GC

Hello World

What's Next